COS 429: Computer Vision, Fall 2018

Homework 1: Color images

Due date: Sep 25, 11:59 PM

Overview

In this homework you will learn to manipulate image in MATLAB.

This is a two part homework. In the first part you will be coloring historic Prokudin-Gorskii images from intensity measurements of red, green and blue color channels. You must first compute an alignment of the three channels to produce a single color image. In the second part you will implement several demosaicing algorithms to produce color images from Bayer pattern images. This is how digital cameras capture color.

Part 1: Coloring Prokudin-Gorskii images

[Sergei Mikhailovich Prokudin-Gorskii](#) was a color photographer far ahead of his time. He undertook a photographic survey of the Russian Empire for Tsar Nicholas II and was able to capture color pictures before color cameras were inveted. His idea was to simply take three pictures of each scene, each with a red, green and blue color filter which could then be used to obtain a color image. There was no way of printing these back in the day, so he envisioned complex display devices to show these. However these were never made, but his pictures survived. In this homework you will reconstruct a color image from his scans of his photographs as seen in Figure 1.

The key step to compute the color image is to align the channels. Alas, since the camera moved between each shot, these channels are slightly displaced from each other. The simplest way is to keep one channel fixed say R, and align the G, B channels to it by searching over displacements in the range [-15, 15] both horizontally and vertically. Pick the display that maximizes similarity between the channels. One such measure is dot product, i.e, `R'G`. Another is normalized cross-correlation, which is simply the dot product between l2 normalized `R, G` vectors.

Figure 1: Example image from the Prokudin-Gorskii collection. Note the colors are in **B, G, R** order from the top to bottom (and not **R, G, B**)!

Before you start aligning the Prokudin-Gorskii images, you will test your code on synthetic images which have been randomly shifted. Your code should correctly discover the inverse of the shift.

Code

Run `evalAlignment.m` on the MATLAB command prompt inside the code directory. This should produce the following output.

Note the actual 'gt shift' might be different since it is randomly generated.

```
Evaluating alignment ..
1 balloon.jpeg
      gt shift: ( 1,11) ( 4, 5)
    pred shift: ( 0, 0) ( 0, 0)
2 cat.jpg
      gt shift: (-5, 5) (-6,-2)
    pred shift: ( 0, 0) ( 0, 0)
 ...
```

The code loads a set of images, randomly shifts the color channels and provides them as input to `alignChannels.m`. Your goal is to implement this function. A correct implementation should obtain the shifts that is the negative of the ground-truth shifts, i.e., the following output:

```
Evaluating alignment ..
1 balloon.jpeg
      gt shift: ( 13, 11)  ( 4, 5)
    pred shift: (-13,-11)  (-4,-5)
2 cat.jpg
      gt shift: (-5, 5) (-6,-2)
    pred shift: ( 5,-5) ( 6, 2)
 ...
```

Once you are done with that, run `alignProkudin.m`. This will call your function to align images from the Prokudin-Gorskii collection. The output is saved to the outDir. Note: if this directory does not exist, you will have to create it first. In your report, show all the aligned images as well as the shifts that were computed by your algorithm.

Tips: Look at functions `circshift()` and `padarray()` to deal with shifting images.

Extra credit

Here are some ideas for extra credit:

- **Boundary effects**: Shifting images causes ugly boundary artifacts. Come up with of a way of avoiding this.
- **Faster alignment**: Searching over displacements can be slow. Think of a way of aligning them faster in a coarse to fine manner. For example you may align the channels by resizing them to half the size and then refining the estimate. This can be done by multiple calls to your `alignChannels()` function.
- **Gradient domain alignment**: Instead of aligning raw channels, you might want to align the edges to avoid overdue emphasis on constant intensity regions. You can do this by first computing an edge intensity image and then using your code to align the channels. Look at the `edge()` function in Matlab.

Part 2: Color image demosaicing

Recall that in digital cameras the red, blue, and green sensors are interlaced in a Bayer pattern. Your goal is to fill the missing values in each channel to obtain a single color image. For this homework you will implement several different interpolation algorithms. The input to the algorithm is a single image `im`, a NxM array of numbers

between 0 and 1. These are measurements in the format shown in Figure 2, i.e., top left `im(1,1)` is red, `im(1,2)` is green, `im(2,1)` is green and `im(2,2)` is blue channel. Your goal is to create a single color image C from these measurements.


Figure 2: Bayer pattern

Start with some non-adaptive algorithms. Implement *nearest neighbour* and *linear* interpolation algorithm. Then implement an adaptive algorithm such as *weighted gradient averaging*. Compare the results. Where are the errors?

Evaluation

We will evaluate the algorithm by comparing your output to the ground truth color image. The input to the algorithm was constructed by artificially mosaicing it. This is not ideal in practice, since there is a demosaicing algorithm that was used to construct the digital image, but we will ignore this for now. We can compute the mean error between each color image and your output and report these numbers for each algorithm.

Code

Run `evalDemosaicing` on the MATLAB command prompt inside the code directory. You should be able to see the output shown in Figure 2. The code loads images from the data directory, artificially mosaics them, and provides them as input to the demosaicing algorithms. The actual code for that is in `demosaicImage.m` file. Right now only the `demosaicImage(im, 'baseline')` is implemented. All the other methods call the baseline methods which is why they produce identical errors. Your job is to implement the following functions in the file:

- `demosaicImage(im, 'nn')` - nearest-neighbour interpolation.
- `demosaicImage(im, 'linear')` - linear interpolation.
- `demosaicImage(im, 'adagrad')` - adaptive gradient interpolation.

The baseline method achieves an average error of 0.1392 across the 10 images in the dataset. Your methods you should be able to achieve substantially better results. As a reference, the linear interpolation method achieves an error of about 0.0174. *You should report the results of all your methods in the form of the same table.*

```
--------------------------------------------------------------------------------
#        image         baseline        nn              linear          adagrad
--------------------------------------------------------------------------------
1        balloon.jpeg   0.179239        0.179239        0.179239        0.179239
2        cat.jpg        0.099966        0.099966        0.099966        0.099966
3        ip.jpg                 0.231587        0.231587        0.231587        0.231587
4        puppy.jpg      0.094093        0.094093        0.094093        0.094093
5        squirrel.jpg   0.121964        0.121964        0.121964        0.121964
6        candy.jpeg     0.206359        0.206359        0.206359        0.206359
7        house.png      0.117667        0.117667        0.117667        0.117667
8        light.png      0.097868        0.097868        0.097868        0.097868
9        sails.png      0.074946        0.074946        0.074946        0.074946
10       tree.jpeg      0.167812        0.167812        0.167812        0.167812
--------------------------------------------------------------------------------
```

```
              average     0.139150        0.139150        0.139150        0.139150
    ---------------------------------------------------------------------------
```

Figure 2: Output of `evalDemosaicing.m`

Tips: You can visualize at the errors by setting the display flag to true in the `runDemosaicing.m`. Avoid loops for speed in MATLAB. Be careful in handling the boundary of the images. It might help to think of various operations as convolutions. Look up MATLAB's `conv2()` function that implements 2D convolutions.

Extra credit

Here are some ideas for extra credit:

- **Transformed color spaces:** Try your previous algorithms by first interpolating the green channel and then transforming the red and blue channels R ← R/G and B ← B/G, i.e., dividing by the green channel and then transforming them back after interpolation. Try other transformations such as logarithm of the ratio, etc (*note: you have to apply the appropriate inverse transform*). Does this result in better images measured in terms of the mean error? Why is this a good/bad idea?
- **Digital Prokudin-Gorskii:** Imagine that Prokudin was equipped with a modern digital camera. Compare the resulting demosaiced images to what he obtained using his method of taking three independent grayscale images. Start with the aligned images of your first part and create a mosaiced image by sampling the channels image using the Bayer pattern. Run your demosaicing algorithms and compare the resulting image with the input color image. How much do you lose? Where are the errors?

Writeup

You are required to turn in a writeup briefly describing your results. Include example images of the output obtained from each part. The writeup should be self contained, describing the key choices you made. If you do any of the suggested extra credit, explain clearly, and include any quantitative or qualitative results you obtained.

Although, this is the primary way in which you will be evaluated you are also required to turn in code for your experiments. For this homework only the functions `alignChannels.m` and `demosaicImage.m` should be changed.

Submission

What to submit? We will be using both gradescope and blackboard for homework submission. Following are the things you need to submit:

- A report for part 1 (and associated extra credits) in Gradescope (PDF format).
- A report for part 2 (and associated extra credits) in Gradescope (PDF format).
- Code for part 1, the file `alignChannels.m`
- Code for part 2, the file `demosaicImage.m`
- Code for extra credits.
- All the codes need to be submitted as a zip file with name "netID.zip", eg. "agoyal.zip" on Blackboard.

Acknowledgements